

# Databases Autumn 2025

## Hand-In Exercise 4

November 22, 2025

Aiysha Frutiger  
Jannick Seper  
Luis Tritschler

Total Points	

Task	Points

## Task 1

a)

```
1 CREATE FUNCTION checkBalance (cid IN NUMBER)
2 RETURN VARCHAR2
3 IS
4     balanceCalculated NUMBER(10,2);
5     balanceSaved NUMBER(10,2);
6 BEGIN
7     SELECT SUM(amount)
8     INTO balanceCalculated
9     FROM BankTransaction
10    WHERE customer_id = cid;
11    SELECT balance
12    INTO balanceSaved
13    FROM Customer
14    WHERE customer_id = cid;
15
16    IF balanceCalculated = balanceSaved THEN
17        RETURN 'Consistent';
18    ELSE
19        RETURN 'Inconsistent';
20    END IF;
21 END;
```

b)

```
1 CREATE PROCEDURE checkAllBalances (consistent OUT VARCHAR2)
2 IS
3     inconsistent BOOLEAN := FALSE;
4     CURSOR c IS SELECT customer_id FROM Customer;
5 BEGIN
6     FOR record IN c LOOP
7         IF check_balance(record.customer_id) =
8             'Inconsistent' THEN
9             inconsistent := TRUE;
10            END IF;
11        END LOOP;
12
13        IF inconsistent THEN
14            consistent := 'Inconsistent';
15        ELSE
16            consistent := 'Consistent';
17        END IF;
18    END;
```

c)

```
1 CREATE PROCEDURE checkReceive (
2     cid                      IN NUMBER,
3     amount                   IN NUMBER(10,2),
4     description              IN VARCHAR2)
5 IS
6     customerExists          NUMBER;
7     customerNotExists       EXCEPTION;
8     maxTransaction          BankTransaction.transaction_id%TYPE;
9 BEGIN
10    SELECT COUNT(*)
11        INTO customerExists
12    FROM Customer
13   WHERE customer_id = cid;
14
15   IF customerExists = 0 THEN RAISE customerNotExists;
16 END IF;
17
18   UPDATE Customer
19   SET balance = balance + amount
20   WHERE customer_id = cid;
21
22   SELECT MAX(transaction_id)
23     INTO maxTransaction
24   FROM BankTransaction;
25
26   INSERT INTO BankTransaction
27   VALUES (maxTransaction + 1, cid, SYSDATE,
28                           amount, description);
29   COMMIT;
30
31 EXCEPTION
32   WHEN customerNotExists THEN
33       ROLLBACK;
34   WHEN OTHERS THEN
35       ROLLBACK;
36 END;
```

## Task 2

a)

We first determine how many tuples fit on one data page and then how many pages are required per relation. The page size is 4096 B, of which 96 B are reserved for the page header. Thus, the usable space per page is

$$4096 \text{ B} - 96 \text{ B} = 4000 \text{ B}.$$

Since data pages are filled only up to 90%, the effective capacity for tuples is

$$0.9 \cdot 4000 \text{ B} = 3600 \text{ B}.$$

Each tuple additionally requires one slot array pointer of 6 B. Hence, the effective record size is

$$\text{record size} = \text{tuple size} + 6 \text{ B}.$$

The number of tuples per page is then

$$\text{tuples per page} = \left\lfloor \frac{3600}{\text{record size}} \right\rfloor,$$

and the number of pages is

$$\text{pages} = \left\lceil \frac{\#\text{tuples}}{\text{tuples per page}} \right\rceil.$$

**Movie** Average tuple size: 100 B, record size:  $100 + 6 = 106 \text{ B}$ .

$$\text{tuples per page} = \left\lfloor \frac{3600}{106} \right\rfloor = 33, \quad \text{pages} = \left\lceil \frac{1000}{33} \right\rceil = 31.$$

**Scene** Average tuple size: 50 B, record size:  $50 + 6 = 56 \text{ B}$ .

$$\text{tuples per page} = \left\lfloor \frac{3600}{56} \right\rfloor = 64, \quad \text{pages} = \left\lceil \frac{10000}{64} \right\rceil = 157.$$

**Person** Average tuple size: 40 B, record size:  $40 + 6 = 46 \text{ B}$ .

$$\text{tuples per page} = \left\lfloor \frac{3600}{46} \right\rfloor = 78, \quad \text{pages} = \left\lceil \frac{6000}{78} \right\rceil = 77.$$

**Actor** Same tuple size as **Person**, hence same record size:  $40 + 6 = 46 \text{ B}$ .

$$\text{tuples per page} = \left\lfloor \frac{3600}{46} \right\rfloor = 78, \quad \text{pages} = \left\lceil \frac{30000}{78} \right\rceil = 385.$$

Relation	Tuples per page	Number of pages
Movie	33	31
Scene	64	157
Person	78	77
Actor	78	385

b)

The page size is 4096 B with a page header of 96 B, thus

$$4096 \text{ B} - 96 \text{ B} = 4000 \text{ B}$$

are available per page. Index pages are filled up to 70%, hence the effective capacity is

$$C = 0.7 \cdot 4000 \text{ B} = 2800 \text{ B}.$$

Pointers in the B<sup>+</sup>-tree and tuple identifiers (TIDs) have a length of 6 B, data type NUMBER has 10 B, and the average MovieTitle length is 90 B.

**MovieIDIdx** Key size  $k = 10 \text{ B}$ . In a leaf node, one entry consists of (key, TID) with size  $10 + 1 \cdot 6 = 16 \text{ B}$ . We reserve 2 pointers of 6 B on each leaf page (e.g. sibling pointers):

$$t_{\text{leaf}} = \left\lfloor \frac{(4096 - 96 - (2 \cdot 6)) \cdot 0.7}{10 + (1 \cdot 6)} \right\rfloor = 174.$$

For inner nodes, one entry consists of (key, child pointer) of size  $10 + 6 = 16 \text{ B}$ , plus an additional pointer of 6 B:

$$e_i = \left\lfloor \frac{(4096 - 96) \cdot 0.7 - 6}{10 + 6} \right\rfloor = 174.$$

With 1000 movies, the number of leaf nodes is

$$n_{\text{leaf}} = \left\lceil \frac{1000}{174} \right\rceil = 6.$$

An inner node can point to  $e_i + 1 = 175$  children, thus a single root node suffices:

$$n_{\text{inner}} = \left\lceil \frac{6}{175} \right\rceil = 1.$$

The total index size is

$$i_{\text{size}} = (6 + 1) \cdot 4 = 28 \text{ KiB}.$$

**MovieTitleIdx** Key size  $k = 90 \text{ B}$ . Leaf entry size:  $90 + 1 \cdot 6 = 96 \text{ B}$ :

$$t_{\text{leaf}} = \left\lfloor \frac{(4096 - 96 - (2 \cdot 6)) \cdot 0.7}{90 + (1 \cdot 6)} \right\rfloor = 29.$$

Inner node:

$$e_i = \left\lfloor \frac{(4096 - 96) \cdot 0.7 - 6}{90 + 6} \right\rfloor = 29.$$

With 1000 movies:

$$n_{\text{leaf}} = \left\lceil \frac{1000}{29} \right\rceil = 35.$$

Each inner node can have up to  $e_i + 1 = 30$  children, so the number of inner nodes one level above the leaves is

$$n_{\text{inner1}} = \left\lceil \frac{35}{30} \right\rceil = 2,$$

and with one root node above:

$$n_{\text{inner}2} = \left\lceil \frac{2}{30} \right\rceil = 1, \quad n_{\text{inner}} = 2 + 1 = 3.$$

Thus, the total index size is

$$i_{\text{size}} = (35 + 3) \cdot 4 = 152 \text{ KiB}.$$

**ActorIdx** There are 30000 tuples in **Actor** and 6000 tuples in **Person**, so on average there are

$$r_k = \frac{30000}{6000} = 5$$

**Actor** entries per PID. In the leaf nodes of **ActorIdx**, we store one entry per distinct key PID, together with a list of  $r_k$  TIDs. Hence, the average leaf entry size is

$$10 + r_k \cdot 6 = 10 + 5 \cdot 6 = 40 \text{ B}.$$

We again reserve 2 pointers of 6 B per leaf page:

$$t_{\text{leaf}} = \left\lceil \frac{(4096 - 96 - (2 \cdot 6)) \cdot 0.7}{10 + (5 \cdot 6)} \right\rceil = 69.$$

For inner nodes, we only store (key, child pointer):

$$e_i = \left\lceil \frac{(4096 - 96) \cdot 0.7 - 6}{10 + 6} \right\rceil = 174.$$

We have 6000 distinct keys (PIPs), hence

$$n_{\text{leaf}} = \left\lceil \frac{6000}{69} \right\rceil = 87.$$

An inner node can point to up to  $e_i + 1 = 175$  children, so

$$n_{\text{inner}} = \left\lceil \frac{87}{175} \right\rceil = 1.$$

The total index size is therefore

$$i_{\text{size}} = (87 + 1) \cdot 4 = 352 \text{ KiB}.$$

Index	Entries per leaf nodes		Number of leaf nodes		Index Size [KiB]
	inner nodes	inner nodes	leaf nodes	inner nodes	
<b>MovieIDIdx</b>	174	174	6	1	28
<b>MovieTitleIdx</b>	29	29	35	3	152
<b>ActorIdx</b>	69	174	87	1	352

c)

```
SELECT * FROM Movie WHERE MovieID = 4711;
SELECT * FROM Movie WHERE MovieTitle = 'Opelgang';
SELECT * FROM Actor WHERE PID = 1199;
```

**Movie.MovieID** For `MovieIDIdx` we have  $n_{leaf} = 6$  leaf nodes and a maximum fanout of 175 children per inner node. Hence, the height of the  $B^+$ -tree (including root and leaves) is

$$h = \lceil \log_{175} 6 \rceil + 1 = 2.$$

Thus we access  $h = 2$  index pages plus one data page in `Movie`:

$$\text{total page accesses} = 2 + 1 = 3.$$

**Movie.MovieTitle** For `MovieTitleIdx` we have  $n_{leaf} = 35$  leaf nodes and a fanout of 30 children per inner node. The height is

$$h = \lceil \log_{30} 35 \rceil + 1 = 3.$$

Again, we access one data page in `Movie`:

$$\text{total page accesses} = 3 + 1 = 4.$$

**Actor.PID** For `ActorIdx` we have  $n_{leaf} = 87$  leaf nodes and a fanout of 175 children per inner node. The height is

$$h = \lceil \log_{175} 87 \rceil + 1 = 2.$$

On average there are  $r_k = 5$  `Actor` tuples per `PID`, so we expect to access 5 data pages in `Actor`:

$$\text{total page accesses} = 2 + 5 = 7.$$

Relation	Attribute	Page accesses
<code>Movie</code>	<code>MovieID</code>	3
<code>Movie</code>	<code>MovieTitle</code>	4
<code>Actor</code>	<code>PID</code>	7